

# Lab 04 – Scheduling



**Instructor:** Lorenzo De Carli, University of Calgary ([lorenzo.decarli@ucalgary.ca](mailto:lorenzo.decarli@ucalgary.ca))  
*Slides by Lorenzo De Carli, based on material by Robert Walls (WPI)*

# In this lab

- Build a **scheduler simulator**
  - Receive as input a **sequence of jobs** in a file
  - **Simulate execution** of those jobs
  - Output an **execution trace** with:
    - Job start time
    - Job end time
    - Policy analysis

# Deliverable

## A scheduler executable

- You will be provided with a **template** (including tests and Makefile)
- Compiler should produce a **scheduler.out** file accepting **3 parameters**:
  - A flag (**0** or **1**) detailing whether or not to perform **policy analysis**
    - ...more details about this later
  - Name of the scheduling policy (**FIFO/SJF/STCF/RR/LT**)
  - Length of each **time-slice** (used for RR; ignored otherwise)
  - Name of input job file (e.g. **jobs.txt**)

# Example

## Running the SJF policy w/ no analysis from jobs.txt

`./scheduler.out 0 SJF 2 jobs.txt`

- You can assume **all parameters** are **always specified**
- Here **0** means “no policy analysis”
- **SJF** means to run the SJF scheduling policy
- **2** means a timeslice of 2 time units (meaningless for SJF)
- **jobs.txt** means to read the list of jobs from the file “jobs.txt”

# Input file format

- Each **workload** is defined in a **workload file**.
- Each line of the workload file represents a **different job** in the workload
- Each line consists of **two comma-separated numbers**:
  - the **arrival time**, and
  - the **total amount of simulated time** that job needs to run.

# Input file format - example

0,1

2,5

3,11

5,4

7,1

25,4

← Each line includes **arrival time** and **duration**

You can assume there are **no extra spaces**,  
and that **every line ends with a newline**

# So... what's about these jobs?

- To be clear, those are **not actual jobs**
- The scheduler should **simulate** that sequence of jobs by appropriately **computing when they start and end**
- There is no need to actually run anything (except the **scheduler calculations**)

# Job list data structure

- The scheduler uses the job file to initialize a **job list data structure**
- In practice, this should be a **linked list**
- Each job should be assigned an **id** based on the **line number in the file**
- The job on the **first line** should be assigned an id of **0**; the job on the **second line** should be assigned an id of **1**; and so on



# Job list data structure /2

This is just an example...

```
struct job {  
    int id;  
    int arrival;  
    int length;  
    // other meta data  
    struct job *next;  
};
```

# Some more things...

- Your scheduler should account for periods when there are **no jobs** to run
  - That is, all the arrived jobs **have completed** before the new jobs arrive
  - In other words, **the CPU can be idle**
- For pre-emptive scheduling policies: also note that a job (or the remaining duration of a job) may last **less than the duration of a time slice**

# Implementing policies

# Implementing FIFO

- The **FIFO policy** is one of the simplest scheduling policies
  - Good starting point! 😊
- The FIFO policy states that **jobs are scheduled in order of their arrival**
- Each job **runs to completion**
- To be clear: there is **no preemption** for this FIFO policy.

# Example scheduler output...

## ...when running FIFO

```
$ ./scheduler.out 0 FIFO 2 tests/3.in
```

Execution trace with FIFO:

t=0: [Job 0] arrived at [0], ran for: [20]

t=20: [Job 1] arrived at [0], ran for: [19]

t=39: [Job 2] arrived at [1], ran for: [18]

t=57: [Job 3] arrived at [1], ran for: [17]

t=74: [Job 4] arrived at [2], ran for: [16]

t=90: [Job 5] arrived at [3], ran for: [15]

t=105: [Job 6] arrived at [4], ran for: [14]

End of execution with FIFO.

# Implementing SJF

(aka “Shortest Job First”)

- **SJF** always picks the job with the **shortest runtime** to run next
- We again assume that a job will run to completion **before the next is started**
- If two jobs need the same amount of time, SJF breaks the tie by favoring **the job that arrived earlier**
- Your SJF scheduler should account for periods when there are **no jobs** to run
  - That is, all the arrived jobs **have completed** before the new jobs arrive
  - In other words, **the CPU can be idle**

# Example scheduler output...

...when running SJF

```
$ ./scheduler 0 SJF 2 tests/8.in
```

Execution trace with SJF:

t=0: [Job 0] arrived at [0], ran for: [1]

t=2: [Job 1] arrived at [2], ran for: [5]

Note that the CPU was  
**idle** for 1 tick here

t=7: [Job 4] arrived at [7], ran for: [1]

t=8: [Job 3] arrived at [5], ran for: [4]

t=12: [Job 2] arrived at [3], ran for: [11]

t=25: [Job 5] arrived at [25], ran for: [4]

End of execution with SJF.

# Implementing STCF

(aka “Shortest Time to Completion”)

- The **STCF policy** is a preemptive version of SJF
- STCF only makes scheduling decisions when jobs arrive/complete
- When a new job arrives, STCF must run find out which jobs has the shortest amount of time left, and run that one
- When a job complete, the same process is also followed



# Example scheduler output...

## ...when running STCF

Execution trace with STCF:

t=100: [Job 0] arrived at [100], ran for: [10]

t=110: [Job 1] arrived at [110], ran for: [10]

t=120: [Job 2] arrived at [110], ran for: [10]

t=130: [Job 0] arrived at [100], ran for: [90]

t=220: [Job 3] arrived at [220], ran for: [20]

t=240: [Job 4] arrived at [220], ran for: [30]

End of execution with STCF.

# Implementing RR

(aka “Round-Robin”)

- **RR** runs each job in turn for the duration of the time slice  $S$
- Note that not all jobs may arrive at the same time!
  - If a scheduling decision is being made at time  $T$ , only jobs arrived at or before  $T$  should be considered
- Once a job has been run for  $S$  ticks, its duration must be diminished by  $S$

# Example scheduler output...

## ...when running RR

Execution trace with RR:

t=0: [Job 0] arrived at [0], ran for: [2]

t=2: [Job 0] arrived at [0], ran for: [2]

t=4: [Job 0] arrived at [0], ran for: [1]

t=5: [Job 1] arrived at [5], ran for: [2]

t=7: [Job 2] arrived at [5], ran for: [2]

t=9: [Job 2] arrived at [5], ran for: [1]

t=14: [Job 3] arrived at [14], ran for: [1]

End of execution with RR.

# Implementing LT

(aka “Lottery Scheduling”)

- Strictly speaking, a **lottery scheduler** does not have to be preemptive...
- ...however, here we are going to implement it as such
- The lottery scheduler assigns a number of **tickets** to each job. Then:
  - **Extract ticket**  $T$  and run the job  $J_T$  to which the  $T$  belongs to for  $S$  ticks
  - Reduce the duration of  $J_T$  by  $S$
  - Look at all jobs and **run the lottery again**

# Extract ticket?

- > **Simple implementation:** Use a **linked list** of jobs and the allotted number of tickets.
- > Extract the winning number using a **random number generator**
- > **Traverse the list** and use a simple **counter** and stop when that counter exceeds the winning number.
- > See lecture slides for more details...

# How do I assign tickets to jobs?

**We are going to keep it simple**

- Assign 100 tickets to the first job that arrives
- 200 tickets to the next
- And so on...

# A note on randomness

- In order to ensure compliance with the test, you need to pre-initialize the random number generator

Look at code template: `srand(42);`

- Later, every time the scheduler makes a decision you can select the winning ticket using: `int winning_ticket = rand() % total_tickets;`
- ... and then, use the linked list approach discussed in class to pick the winning process
- We'll be somewhat flexible w/ compliance with the tests for the LT part

# Example scheduler output...

## ...when running LT

Execution trace with LT:

```
t=0: [Job 2] arrived at [0], ran for: [10]
t=10: [Job 1] arrived at [0], ran for: [10]
t=20: [Job 0] arrived at [0], ran for: [10]
t=30: [Job 2] arrived at [0], ran for: [10]
t=40: [Job 2] arrived at [0], ran for: [10]
t=50: [Job 1] arrived at [0], ran for: [10]
t=60: [Job 2] arrived at [0], ran for: [10]
t=70: [Job 1] arrived at [0], ran for: [10]
t=80: [Job 2] arrived at [0], ran for: [10]
t=90: [Job 2] arrived at [0], ran for: [10]
t=100: [Job 2] arrived at [0], ran for: [10]
```

End of execution with LT.



# Policy analysis

# Policy analysis?

- In this part of this project, you will add code to the scheduler to help it evaluate the **performance** of the previously implemented **policies**
- Your code will measure **three metrics**:
  - Response time
  - Turnaround time
  - Wait time

# Metric definitions

- Assume:
  - $T_a$  is the job **arrival time**
  - $T_s$  is the job **start time**
  - $T_c$  is the job **completion time**
- Then:
  - **Response time** is  $T_s - T_a$
  - **Turnaround time** is  $T_c - T_a$

# Wait time

- Wait time is the total accumulated amount of time a job spends waiting while other jobs run
- If a job arrives at 0, starts at **6**, runs for 2 ticks, then wait for **4** ticks, then run for 2 ticks, is overall wait time is  **$6 + 4 = 10$**
- Note: for non-pre-emptive policies, wait time and response time are the same!

# Scheduler policy analysis

- The modified scheduler should output, for each metric:
  - The **per-job value** of the metric
  - The **average value** of the metric **across all jobs**

# Example FIFO scheduler output...

## ...with metrics

```
$ ./scheduler 1 FIFO 2 tests/3.in
```

```
Execution trace with FIFO:
```

```
t=0: [Job 0] arrived at [0], ran for: [20]
```

```
t=20: [Job 1] arrived at [0], ran for: [19]
```

```
t=39: [Job 2] arrived at [1], ran for: [18]
```

```
t=57: [Job 3] arrived at [1], ran for: [17]
```

```
t=74: [Job 4] arrived at [2], ran for: [16]
```

```
t=90: [Job 5] arrived at [3], ran for: [15]
```

```
t=105: [Job 6] arrived at [4], ran for: [14]
```

```
End of execution with FIFO.
```

```
Begin analyzing FIFO:
```

```
Job 0 -- Response time: 0 Turnaround: 20 Wait: 0
```

```
Job 1 -- Response time: 20 Turnaround: 39 Wait: 20
```

```
Job 2 -- Response time: 38 Turnaround: 56 Wait: 38
```

```
Job 3 -- Response time: 56 Turnaround: 73 Wait: 56
```

```
Job 4 -- Response time: 72 Turnaround: 88 Wait: 72
```

```
Job 5 -- Response time: 87 Turnaround: 102 Wait: 87
```

```
Job 6 -- Response time: 101 Turnaround: 115 Wait: 101
```

```
Average -- Response: 53.43 Turnaround 70.43 Wait 53.43
```

```
End analyzing FIFO.
```

# Example SJF scheduler output...

## ...with metrics

```
$ ./scheduler 1 SJF 2 tests/8.in
```

```
Execution trace with SJF:
```

```
t=0: [Job 0] arrived at [0], ran for: [1]
```

```
t=2: [Job 1] arrived at [2], ran for: [5]
```

```
t=7: [Job 4] arrived at [7], ran for: [1]
```

```
t=8: [Job 3] arrived at [5], ran for: [4]
```

```
t=12: [Job 2] arrived at [3], ran for: [11]
```

```
t=25: [Job 5] arrived at [25], ran for: [4]
```

```
End of execution with SJF.
```

```
Begin analyzing SJF:
```

```
Job 0 -- Response time: 0 Turnaround: 1 Wait: 0
```

```
Job 1 -- Response time: 0 Turnaround: 5 Wait: 0
```

```
Job 2 -- Response time: 9 Turnaround: 20 Wait: 9
```

```
Job 3 -- Response time: 3 Turnaround: 7 Wait: 3
```

```
Job 4 -- Response time: 0 Turnaround: 1 Wait: 0
```

```
Job 5 -- Response time: 0 Turnaround: 4 Wait: 0
```

```
Average -- Response: 2.00 Turnaround 6.33 Wait 2.00
```

```
End analyzing SJF.
```

# Example RR scheduler output...

## ...with metrics

Execution trace with RR:

t=0: [Job 0] arrived at [0], ran for: [1]

t=1: [Job 1] arrived at [1], ran for: [2]

t=3: [Job 2] arrived at [3], ran for: [2]

t=5: [Job 1] arrived at [1], ran for: [2]

t=7: [Job 2] arrived at [3], ran for: [2]

t=9: [Job 1] arrived at [1], ran for: [1]

End of execution with RR.

Begin analyzing RR:

Job 0 -- Response time: 0    Turnaround: 1    Wait: 0

Job 1 -- Response time: 0    Turnaround: 9    Wait: 4

Job 2 -- Response time: 0    Turnaround: 6    Wait: 2

Average -- Response: 0.00    Turnaround 5.33    Wait 2.00

End analyzing RR.



# Code template

- **scheduler.c**: template file to complete to implement the scheduler
- To complete the lab, you must extend the template to implement the required functionality
- As in lab 3, you must only submit the *scheduler.c* file, nothing else!

# There are tests!

## Look in the code template folder

- `test_fifo/`: FIFO policy tests
- `test_sjf/`: SJF policy tests
- `test_stcf/`: STCF policy tests
- `test_rr/`: RR policy tests
- `test_lt/`: LT policy tests
- Each test folder includes tests for runs with and without analysis
- We will also use these tests for grading

# How to run tests?

- You can run individual tests by doing (e.g., for feature 1):

```
cd test_fifo
```

```
python3 test_fifo.py
```

- From the output, it will be clear which tests passed and which ones do not
- Look into the **.in** files in each test directory to understand what specifically is being tested (the corresponding expected output is in the **.out** files)
- To run all tests, run **python3 test\_all.py**
- Note, the tests assume that there exists a **scheduler.out** executable in the main project directory (remember to run **make** before running the tests!)

# Grading rubric

...the part everyone cares about!

- You'll have until **11:59PM of the day before the next lab** to upload the **complete solution**. That will be graded as follows:
  - Correct **FIFO implementation**: 2 pts
  - Correct **SJF implementation**: 2 pts
  - Correct **STCF analysis**: 2 pts
  - Correct **RR analysis**: 2 pts
  - Correct **LT analysis**: 2pts

# Submission instructions

- You need to submit **a single file to the D2L “Lab 3” dropbox**
  - You have until **11:59 PM on the day before the next lab** to submit your completed lab solution, in a single *scheduler.c* file
  - **Please do not compress the file!**
- **THE SUBMISSION MUST ONLY CONSIST OF scheduler.c, NOTHING ELSE**
  - NO tests
  - NO readme
- This submission will be graded **10 Pts** according to the rubric

**That's all!**